



Wheel Bandit Benchmark (Osband et al., NeurIPS 2018 / Snoek et al., ICLR 2018)

Deep Bayesian Exploration & Policy Recovery



(a) Ground Truth Optimal Policy
(Hub=Arm 0, Rim=Jackpot Arms 1-4)




(b) Deep Ensembles Policy
(Trapped in Arm 0; Misses Jackpots)



(c) DP-BNN (Ours) Policy
(Discovers All 4 Jackpot Petals)



(d) Cumulative Regret over Time (Lower is Better)
[Linear Regret = Stuck in Trap, Sub-Linear = Exploration Solved]



(e) Final Total Regret Benchmark

